

Quality Assurance & AI Governance

AI Output Trust Score

Evaluation Framework

A practical, repeatable scoring framework for assessing AI-generated outputs produced through the modular prompting architecture (Variables · Module 1 Research Standard · Module 2 Methodology · Module 3 Output Structure).

Document Control	Detail
Document title	AI Output Trust Score Evaluation Framework
Version	1.0
Audience	AI governance, QA reviewers, prompt engineers, compliance, business stakeholders
Source basis	Modular Prompting Guide (supplied PDF, Pages 1–2)
Total score	100 points across five weighted categories
Review cadence	Quarterly recalibration (recommended assumption — see Section 10)

1. Executive Summary

The AI Output Trust Score is a single, defensible number — out of 100 — that tells a reviewer whether an AI-generated analytical output can be trusted, used, escalated, or rejected. It exists because the difference between a usable AI deliverable and an unusable one is no longer about fluency; it is about whether the output is grounded in evidence, follows the agreed methodology, conforms to the agreed structure, and exposes its own assumptions and uncertainties.

This framework is purpose-built to evaluate outputs produced through the modular prompting architecture described in the supplied source PDF. That architecture separates what changes (Variables) from what is fixed (Module 1 Research Standard, Module 2 Methodology, Module 3 Output Structure). Because the prompt itself enforces a quality bar, a methodology, and a deliverable format, evaluators can — and must — score the resulting output against those exact same anchors.

The framework converts the four failures that the PDF identifies in single-shot prompting (no quality standard, no methodology, no output structure, no data anchoring) into five weighted scoring categories that together total 100 points. Reviewers apply a 0–5 rubric within each category, multiply by the category weight, and sum to a Trust Score. The Trust Score then maps to one of five interpretation bands ranging from Very High Trust (90–100) to Very Low Trust (0–39), each with an associated usage decision and remediation pathway.

Why this framework exists

Single-shot prompts produce outputs that look polished but cannot be trusted, compared, or repeated. The modular prompting framework fixes the prompt side. This Trust Score framework fixes the evaluation side — so that the same output, reviewed by two different people, lands at substantially the same score.

2. Purpose and Scope

This framework applies to AI-generated analytical outputs that were produced using a modular prompt of the form described in the source PDF — that is, a prompt composed of an editable Variables block plus three fixed modules (Research Standard, Methodology, Output Structure).

In scope

- Financial-analysis outputs of the kinds illustrated in the PDF: Market Snapshot, Commodity Review, Financial Health, Portfolio Checkup, and Macro Backdrop.
- Any other analytical deliverable generated through the same modular prompt template, where Variables, Research Standard, Methodology, and Output Structure can be identified.
- Outputs intended for downstream business use — investment memos, internal review packs, client-facing summaries, or analyst handoffs.

Out of scope (and why)

- Free-form chat replies, brainstorming, or exploratory prompting where no Variables/Modules contract exists. The Structural and Modular Framework Alignment category cannot be scored against an absent template.
- Code generation, translation, and other non-analytical tasks. Reviewers should use a task-appropriate rubric for those.
- Confidential model-internals review (e.g., chain-of-thought audits). This framework scores the visible output, not the model's hidden reasoning.

Who uses this framework

- Quality assurance reviewers and AI governance teams, as the primary scorers.
- Prompt engineers, as a feedback loop into module design.
- Compliance reviewers, who rely on the score band to decide whether a human sign-off is required.
- Business stakeholders and analysts consuming the output, as a single, transparent quality signal.

3. Evaluation Principles

Every score awarded under this framework should follow the seven principles below. They exist to make scores defensible, comparable, and resistant to reviewer drift.

Principle	What it means in practice
Objectivity	Score the output on the page, not the reputation of the model, the analyst, or the company under analysis. A confident tone is not evidence; a citation is.
Traceability	Every point awarded or withheld must be tied to a specific observation in the output (a sentence, a table cell, a missing section). Reviewer notes must record where the evidence was — or was not — found.
Prompt alignment	The prompt is the contract. Score adherence to the Variables, the Research Standard, the Methodology, and the Output Structure as supplied — not to an idealised prompt the reviewer wishes had been used.
Evidence-based scoring	Claims without sourced evidence are penalised in Category 3. Claims with sourced evidence that does not actually support them are penalised more heavily — fabricated sourcing is worse than no sourcing.

Principle	What it means in practice
Transparency	Reviewers must show their working: rubric level chosen, raw category scores, weighted contribution, and final total. A Trust Score with no audit trail is itself a red flag.
Currency	Outputs are scored against the data recency expectations stated in the prompt's Coverage Period and Data Sources fields. Stale data presented as current is a Category 5 failure.
Proportional judgment	Minor formatting deviations should not collapse a score; material methodology gaps or fabricated facts should. Category weights and the 0–5 rubric exist to enforce this.

4. Overview of the Trust Score Model

The Trust Score is composed of five weighted categories. Each category is scored on a 0–5 rubric (described in Section 5), and the rubric level is converted into the category's weighted point contribution. The five categories together total exactly 100 points.

Category weights and PDF anchors

#	Category	Max points	Anchored to (from the source PDF)
1	Prompt Adherence and Instruction Compliance	20	Variables block (Company, Identifier, Coverage Period, Peer Set, Purpose & Audience, Data Sources, Accounting/Currency Notes)
2	Structural and Modular Framework Alignment	20	Module 3 — Output Structure (Executive Snapshot → Scorecard → Deep Dive → Peer Context → Risks → Catalysts → Bottom Line + Watchlist) and the 9-step sequence in Module 2
3	Accuracy, Evidence, and Factual Reliability	25	Module 1 — Research Standard: evidence grounding, source citation, calculation transparency. Addresses the PDF's most-flagged failure: 'No Data Anchoring'.

#	Category	Max points	Anchored to (from the source PDF)
4	Reasoning, Judgment, and Evaluation Quality	20	Module 2 — Methodology (Steps 2–9: Trend, Liquidity, Leverage, Cash flow, Efficiency, Peer context, Risk & catalyst, Synthesis)
5	Currency, Risk Awareness, Restrictions, and Transparency	15	Module 1 — Research Standard: data recency checks, explicit assumptions, calibrated language, red-flag surfacing
	Total Trust Score	100	Five categories, 0–5 rubric each, weighted to a single number out of 100

How a category score is calculated

Each category uses a 0–5 rubric. The reviewer's chosen rubric level is converted to the category's weighted contribution as follows:

Rubric level	Multiplier	Plain-English label	What it indicates
5	$1.00 \times \text{max}$	Excellent	Fully meets the criterion with no observable gaps.
4	$0.85 \times \text{max}$	Strong	Meets the criterion with minor, non-material gaps.
3	$0.65 \times \text{max}$	Adequate	Meets most of the criterion but has visible weaknesses that require reviewer notes.
2	$0.40 \times \text{max}$	Weak	Substantial gaps that materially affect the trustworthiness of the output.
1	$0.20 \times \text{max}$	Poor	The criterion is barely addressed; the output is at high risk of being misleading.
0	$0.00 \times \text{max}$	Fail	The criterion is absent, fabricated, or actively harmful (e.g., hallucinated sources).

Worked formula: Trust Score = Σ (rubric multiplier $\times$ category max) across all five categories. The result is rounded to the nearest whole number.

5. Detailed Scoring Framework

Each of the five subsections below specifies, for one category: the score weighting, the criteria, the 0–5 rubric, three illustrative examples, the reviewer checklist, common failure modes, red flags, and consistency guidance. Reviewers should work through them in order; Section 7 provides a single consolidated checklist for use during review.

5.1 Category 1: Prompt Adherence and Instruction Compliance

Field	Detail
Maximum score	20 points
Purpose	Measures the degree to which the AI output respects the seven editable Variables and the explicit instructions supplied in the prompt (Company, Identifier, Coverage Period, Peer Set, Purpose & Audience, Data Sources, Accounting/Currency Notes).
Why it matters for trust	The Variables block is the only part of the prompt that defines what the user actually asked for. If the output drifts from the Variables — wrong company, wrong period, wrong peer set, wrong audience register — every subsequent strength is wasted. Adherence is the foundation of every other score.

Detailed scoring criteria

- The output addresses the exact Company and Identifier supplied; no substitution, no implicit broadening.
- The Coverage Period is respected — no smuggling-in of out-of-scope quarters or years.
- The Peer Set used in comparisons matches the supplied list; no silent additions or omissions.
- The Purpose & Audience drives tone, depth, and emphasis (e.g., a CFO audience is treated differently from a retail-investor audience).
- Only the supplied Data Sources are used; if the output draws on others, it must say so explicitly.
- Accounting/Currency Notes (e.g., USD vs. local currency, IFRS vs. US GAAP) are observed and applied throughout.

- Any restrictions or exclusions stated in the prompt (e.g., 'do not include forward guidance') are honoured.

0–5 scoring rubric

Level	Label	Description (what the reviewer should observe)
5	Excellent	All seven Variables and any prompt restrictions are visibly respected. The output reads as if the prompt template were a hard contract.
4	Strong	All Variables respected; one minor instruction (e.g., audience register) is slightly off but causes no material distortion.
3	Adequate	Five to six Variables respected; one moderate deviation (e.g., wrong Coverage Period for a sub-section, or two extra peers).
2	Weak	Multiple Variables ignored or interpreted loosely; the output partially answers a different question than the one asked.
1	Poor	Most Variables ignored; the output uses a generic template rather than the supplied parameters.
0	Fail	Output is for the wrong company, wrong period, or wrong audience; or it openly violates an explicit restriction.

Examples

High-scoring example (5)

An output that opens by re-stating the seven Variables, names the exact Coverage Period in every section heading, uses precisely the named peer set in the Peer Context section, expresses every monetary figure in the currency specified in the Accounting/Currency Notes, and adopts the language register implied by the Purpose & Audience field.

Medium-scoring example (≈3)

An output that uses the right Company and Coverage Period but adds one peer that was not on the supplied Peer Set, and reports a few figures in USD when the prompt specified GBP — without flagging the conversion.

Low-scoring example (≤1)

An output that analyses the parent group instead of the specified subsidiary, uses trailing-twelve-months data when the prompt asked for FY2024, and addresses a generic 'investor' audience when the prompt specified an internal credit committee.

Reviewer checklist questions

- Are all seven Variables explicitly observable in the output?
- Is the Coverage Period visibly enforced section-by-section?
- Does the Peer Set in the Peer Context section match the supplied list exactly?
- Does the tone and depth match the Purpose & Audience field?
- Is the currency and accounting basis consistent throughout?
- Are any prompt-level restrictions (exclusions, scope limits) honoured?
- Where the output departs from the prompt, is the departure flagged and justified?

Common failure modes

- Silent substitution of an entity (e.g., parent for subsidiary) without disclosure.
- Period drift — switching from the stated Coverage Period to a more convenient one inside the analysis.
- Peer-set inflation, where extra comparable companies are added 'for completeness'.
- Audience mismatch, producing executive-summary tone for a technical credit-committee request, or vice versa.

Red flags that should reduce the score

- Output ignores a stated restriction (e.g., 'do not include forecasts') and includes the restricted content anyway.
- Currency or unit changes mid-document without an explicit note.
- Generic boilerplate that does not reference the supplied Variables at all.

How to apply this score consistently

Score Category 1 first, before reading the body of the output for analysis quality. Cross-check each Variable against the relevant section of the output and tick it off explicitly in your reviewer notes. If three or more Variables are not visibly respected, the category cannot score above 3 regardless of the analytical quality elsewhere.

5.2 Category 2: Structural and Modular Framework Alignment

Field	Detail
Maximum score	20 points
Purpose	Measures whether the output conforms to Module 3's fixed deliverable format (Executive Snapshot → Key Metrics Scorecard → Deep Dive by Pillar → Peer Context → Risks & Red Flags → Catalysts → Bottom Line + Watchlist) and whether the analytical journey follows Module 2's nine-step methodology.
Why it matters for trust	Structural alignment is what makes outputs comparable across companies, periods, and reviewers. The PDF identifies 'No Output Structure' as one of the four single-shot failures: without it, every response looks different and review becomes painful. A score in this category is, in effect, a measure of whether the output is fit for downstream comparison and handoff.

Detailed scoring criteria

- All seven Module 3 sections are present, in the correct order, with appropriate headings.
- The Key Metrics Scorecard is rendered as an actual table — not a paragraph or a bulleted list.
- The Deep Dive by Pillar covers Methodology Steps 3–6 (Liquidity & working capital, Leverage & solvency, Cash flow quality, Efficiency & returns) at minimum.
- Peer Context (Step 7) is a distinct section, not a sentence buried inside the Deep Dive.
- Risks & Red Flags and Catalysts (Step 8) are separated, not merged into a single 'Risks/Catalysts' bucket.
- The Bottom Line + Watchlist (Step 9, Synthesis) provides a clear summary plus forward-looking watchlist items.
- Section depth is balanced — no single section is starved of content while another is over-padded.

0–5 scoring rubric

Level	Label	Description (what the reviewer should observe)
5	Excellent	All seven Module 3 sections present, correctly ordered, appropriately weighted; Module 2 steps are visibly traceable inside the Deep Dive.
4	Strong	All seven sections present; one section is light or out of order, but no information is lost.

Level	Label	Description (what the reviewer should observe)
3	Adequate	Five to six sections present; one section missing or merged. Scorecard may be a list rather than a table.
2	Weak	Three to four sections present; structure improvised; comparability with sibling outputs is reduced.
1	Poor	Output is largely free-form prose with section headings used decoratively rather than structurally.
0	Fail	No discernible alignment with Module 3; output looks like a single-shot reply.

Examples

High-scoring example (5)

An output with the seven Module 3 headings exactly as named, a Scorecard rendered as a four-column table (metric, value, peer median, signal), Deep Dive subsections labelled Liquidity / Leverage / Cash flow / Efficiency, and a Bottom Line + Watchlist section that ends with three named items to monitor next quarter.

Medium-scoring example (≈3)

An output that has six of the seven sections, merges Risks and Catalysts into one section, and presents the Scorecard as a bullet list rather than a table — but otherwise follows the order.

Low-scoring example (≤1)

An output structured as 'Overview / Analysis / Conclusion' with no Scorecard, no Peer Context section, and risks mentioned only in passing inside the conclusion.

Reviewer checklist questions

- Are all seven Module 3 sections present, in order?
- Is the Key Metrics Scorecard a real table?
- Is each of Methodology Steps 3–6 (Liquidity, Leverage, Cash flow, Efficiency) visibly addressed in the Deep Dive?

- Is Peer Context a standalone section?
- Are Risks & Red Flags and Catalysts kept separate?
- Does the Bottom Line + Watchlist contain forward-looking items?
- Is the depth balanced across sections?

Common failure modes

- Decorative section headings that do not match Module 3 (e.g., 'Strengths and Weaknesses' instead of the named sections).
- Scorecard rendered as prose, removing the comparability the table format provides.
- Risks and Catalysts merged, losing the analytical separation between downside and upside drivers.
- Empty Watchlist — the Bottom Line ends without forward-looking monitoring items.

Red flags that should reduce the score

- Two or more Module 3 sections missing entirely.
- Sections present but in the wrong order (e.g., Bottom Line before Deep Dive).
- Methodology Steps 3–6 are not traceable in the Deep Dive — i.e., one or more pillars are absent.

How to apply this score consistently

Score Category 2 by checklist, not impression. Tick each of the seven Module 3 sections, and tick each of the nine Module 2 steps as you find them in the body. The weight is set so that a perfectly-formatted but shallow output cannot win on structure alone; conversely, a structurally weak output cannot be rescued by a strong synthesis.

5.3 Category 3: Accuracy, Evidence, and Factual Reliability

Field	Detail
Maximum score	25 points
Purpose	Measures the degree to which factual claims, numbers, and citations in the output are correct, sourced, and reproducible. Anchored to Module 1 — Research Standard: evidence grounding, source citation, and calculation transparency.
Why it matters for trust	This is the highest-weighted category because the source PDF identifies 'No Data Anchoring' as the failure with the most damaging consequence: hallucination risk. An output that scores well everywhere else but contains fabricated numbers or invented sources is more

Field	Detail
	dangerous, not less, because its surface polish increases the chance the user will act on it without checking.

Detailed scoring criteria

- Every quantitative claim is sourced — the reader can identify the document, page, or dataset behind each number.
- Cited sources actually exist and actually contain the cited data (no fabricated URLs, ticker symbols, or report names).
- Calculations are transparent: ratios show numerator and denominator, growth rates show base and end period, percentages show their derivation.
- Numbers reconcile across sections (e.g., the figure used in the Scorecard matches the figure used in the Deep Dive).
- Where data is estimated or interpolated, this is stated explicitly and the method is described.
- Statements of fact about the company, peers, or macro environment are accurate as of the Coverage Period.
- No claim is materially overstated, understated, or selectively framed.

0–5 scoring rubric

Level	Label	Description (what the reviewer should observe)
5	Excellent	Every claim sourced; sources verifiable; calculations transparent; no factual errors detected.
4	Strong	All material claims sourced; one minor citation gap or one trivial computational rounding error; no fabrication.
3	Adequate	Most claims sourced; one or two unsourced quantitative statements; calculations partially shown; no fabricated sources.
2	Weak	Frequent unsourced claims; calculations not shown; one factual error of moderate consequence.
1	Poor	Most claims unsourced; multiple factual errors; numbers do not reconcile across sections.

Level	Label	Description (what the reviewer should observe)
0	Fail	Hallucinated source(s) detected, or a material number is fabricated, or a factual claim would mislead a decision-maker.

Examples

High-scoring example (5)

An output that, for every figure in the Scorecard, gives the source document and page; shows ratio derivations inline (e.g., 'Current ratio = \$4.2bn / \$2.8bn = 1.50'); reconciles to the Deep Dive figures; and explicitly flags one estimated input as 'Q4 inventory not yet filed; estimated at run-rate'.

Medium-scoring example (≈3)

An output where Scorecard numbers are correct and broadly sourced ('per FY2024 10-K') but two ratios are stated as outputs without showing inputs, and one peer figure is referenced without a specific source.

Low-scoring example (≤1)

An output that quotes a 'Bloomberg consensus' figure with no date, includes a revenue number that does not match the cited 10-K, and contains a paragraph of analysis whose key statistic cannot be traced to any source — or worse, cites a report that does not exist.

Reviewer checklist questions

- Is every number in the Scorecard sourced and verifiable?
- Do the cited sources actually exist (spot-check at least three)?
- Are ratios and growth rates shown with their derivations?
- Do figures reconcile across the Scorecard, Deep Dive, and Bottom Line?
- Are estimates explicitly labelled as such, with method disclosed?
- Are factual statements about the company correct as of the Coverage Period?
- Is any claim selectively framed in a way that misleads?

Common failure modes

- Hallucinated citations — sources that look plausible but do not exist.

- 'Phantom precision' — figures quoted to two decimal places with no underlying data.
- Calculation opacity — final ratio given without numerator/denominator.
- Reconciliation breaks — different figures for the same metric in different sections.
- Selective sourcing — citing only the data points that support the conclusion.

Red flags that should reduce the score

- Any fabricated source, dataset, or report name (immediate score-zero trigger for the category).
- Material number that cannot be reconciled to a cited source.
- Use of phrases like 'industry estimates' or 'market consensus' with no specific reference.
- Ratios presented without inputs.

How to apply this score consistently

Spot-check at least three citations per output. Pick one from the Scorecard, one from the Deep Dive, and one from the Risks section. If any of the three is fabricated, the category cannot score above 1 regardless of the other criteria. This is the most expensive category to under-review; budget time accordingly.

5.4 Category 4: Reasoning, Judgment, and Evaluation Quality

Field	Detail
Maximum score	20 points
Purpose	Measures the analytical quality of the output: whether the reasoning is sound, the synthesis is coherent, the peer context is interpreted (not just listed), and the judgments are calibrated. Anchored to Methodology Steps 2–9 (Trend, Liquidity, Leverage, Cash flow, Efficiency, Peer context, Risk & catalyst map, Synthesis).
Why it matters for trust	An output can be perfectly compliant, perfectly structured, and perfectly sourced and still fail to add value if its reasoning is shallow. This category captures the difference between an AI output that summarises and one that interprets.

Detailed scoring criteria

- Trend analysis explains direction, magnitude, and likely drivers — not just the change in numbers.
- Each Deep Dive pillar reaches a conclusion, not merely a description.

- Peer Context interprets relative position (where does the company stand and why), not just lists peer values.
- The Risk & Catalyst map links each risk and catalyst to a plausible mechanism and a plausible magnitude.
- The Synthesis reconciles tensions between pillars (e.g., strong margins but deteriorating cash conversion) rather than averaging them away.
- Language is calibrated — degrees of confidence are explicit ('likely', 'possible', 'highly uncertain') and used appropriately.
- The output identifies what the analyst should watch next, not just what has happened.

0–5 scoring rubric

Level	Label	Description (what the reviewer should observe)
5	Excellent	Every pillar reaches a defended conclusion; peer context is interpreted; synthesis reconciles tensions; language is calibrated; watchlist is specific.
4	Strong	Reasoning is sound across pillars; one pillar is descriptive rather than evaluative; synthesis is coherent.
3	Adequate	Reasoning is acceptable in two-to-three pillars; one or two pillars are descriptive; peer context is listed rather than interpreted.
2	Weak	Most pillars are descriptive; synthesis is a recap; few defended conclusions.
1	Poor	Output is a paraphrase of the inputs with no analytical value added.
0	Fail	Conclusions contradict the evidence, or the synthesis is internally inconsistent in a way that would mislead a decision-maker.

Examples

High-scoring example (5)

An output that, in the Cash flow pillar, observes an FCF/net income ratio falling below 0.7 for two consecutive years, traces it to extending DSO and rising inventory, links it to the Liquidity pillar, and concludes — calibrated as 'likely, not certain' — that working-capital pressure rather than profitability is the main concern.

Medium-scoring example (≈3)

An output that lists the same FCF/net income figures, notes they have declined, but does not connect the decline to any specific mechanism or to other pillars; the synthesis simply restates each pillar in turn.

Low-scoring example (≤1)

An output that describes each pillar in isolation, makes no comparative interpretation against peers, uses uniformly confident language ('the company is strong / weak'), and concludes with a generic recommendation that does not follow from the analysis.

Reviewer checklist questions

- Does each Deep Dive pillar reach a defended conclusion?
- Is Peer Context interpreted, not just listed?
- Are risks and catalysts linked to mechanism and magnitude?
- Does the Synthesis reconcile tensions between pillars?
- Is language calibrated — are degrees of confidence explicit and appropriate?
- Does the Watchlist contain forward-looking items, not retrospective summaries?
- Would two analysts reading this output reach broadly the same takeaway?

Common failure modes

- Description without interpretation — peer values listed but not explained.
- Pillar isolation — pillars analysed independently with no synthesis across them.
- Confidence creep — 'will' used where 'may' is supported; or, conversely, hedging used to dodge a clearly supported conclusion.
- Watchlist as recap — backward-looking summary instead of forward-looking monitoring items.

Red flags that should reduce the score

- Internal contradiction — synthesis conclusion not supported by, or in tension with, the pillar conclusions.
- Use of recommendation language ('buy', 'sell', 'hold') without any supporting evidence chain.
- Pillars given equal weight when the evidence clearly elevates one pillar's importance.

How to apply this score consistently

Read the Synthesis first, then back-solve: does each conclusion in the Synthesis trace to a defended pillar conclusion above it? Where the Synthesis introduces a claim that is not earned in the body, deduct from this category, not Category 3 — this is a reasoning failure, not a fact failure.

5.5 Category 5: Currency, Risk Awareness, Restrictions, and Transparency

Field	Detail
Maximum score	15 points
Purpose	Measures whether the output handles time, uncertainty, and prompt-level restrictions transparently. Anchored to Module 1 — Research Standard: data recency checks, explicit assumptions, calibrated language, and red-flag surfacing.
Why it matters for trust	Even a fully accurate output becomes untrustworthy if it presents stale data as current, hides its assumptions, ignores stated restrictions, or fails to surface red flags. This category captures the 'meta-quality' of the output — its honesty about its own limitations.

Detailed scoring criteria

- Data recency is explicit — every key data point is dated.
- Where data is older than the Coverage Period would imply, the staleness is flagged.
- All material assumptions are stated up front, not buried.
- Red flags identified during analysis are surfaced in the Risks & Red Flags section, not soft-pedalled in passing.
- Restrictions stated in the prompt (e.g., 'do not provide investment advice') are honoured visibly.
- Limitations of the analysis (data gaps, peer comparability issues, accounting differences) are disclosed.
- Forward-looking statements use calibrated, conditional language rather than asserting the future as fact.

0–5 scoring rubric

Level	Label	Description (what the reviewer should observe)
5	Excellent	Every data point dated; all assumptions explicit; red flags surfaced; restrictions honoured; limitations disclosed.

Level	Label	Description (what the reviewer should observe)
4	Strong	Most data dated; major assumptions stated; one minor limitation not disclosed; restrictions honoured.
3	Adequate	Some data undated; some assumptions implicit; red flags mentioned but understated; one minor restriction softly breached.
2	Weak	Stale data presented as current in places; assumptions largely implicit; red flags missing or buried.
1	Poor	No date stamping; no assumption disclosure; restrictions ignored; output reads as if uncertainty did not exist.
0	Fail	Stale data presented as current with material consequence, or a stated restriction is openly violated, or red flags identifiable in the data are actively suppressed.

Examples

High-scoring example (5)

An output that opens its Scorecard with 'all figures as of Q3 2024 unless dated otherwise', flags one data point ('Q4 2024 inventory not yet filed — Q3 figure used'), states two assumptions explicitly, surfaces a deteriorating-receivables red flag in the Risks section, and uses 'may' / 'likely' / 'is highly uncertain' language consistently in the Catalysts section.

Medium-scoring example (≈3)

An output that dates the headline figures but not the peer figures, states one of three relevant assumptions, and mentions a red flag only in passing inside the Cash flow pillar without elevating it to the Risks section.

Low-scoring example (≤1)

An output that uses 'currently' and 'as of today' without dates, makes no assumptions explicit, presents trailing data as current, ignores a clearly stated 'no forward guidance' restriction, and uses uniformly confident language about the future.

Reviewer checklist questions

- Is every key data point dated?
- Where data is older than the Coverage Period, is the staleness flagged?
- Are material assumptions stated up front?
- Are red flags surfaced in the Risks section, not buried in pillars?
- Are all stated prompt restrictions visibly honoured?
- Are analytical limitations (data gaps, peer comparability) disclosed?
- Is forward-looking language calibrated and conditional?

Common failure modes

- Date-free framing — 'currently' or 'as of today' without an actual date.
- Hidden assumptions — key inputs (e.g., FX rate, growth assumption) not stated.
- Red-flag suppression — issues acknowledged inside the analysis but not elevated to Risks.
- Restriction creep — output drifts past stated restrictions section by section.

Red flags that should reduce the score

- Output presents data older than the Coverage Period as if it were current, with no flag.
- An explicit prompt restriction is breached.
- Forward-looking statements asserted as fact ('will', 'definitely') without supporting evidence.

How to apply this score consistently

Score this category last, after the analytical categories. Many of its failures are visible only once the reviewer knows what the analysis actually concluded, for example, a missing red flag is recognisable only when the underlying data has been checked. Use the Trust Score's overall direction as a sanity check: a strong analytical output with weak transparency often deserves a band downgrade rather than a small deduction here.

6. Trust Score Interpretation Bands

The Trust Score (0–100) maps to one of five bands. The band determines whether the output can be used as-is, requires human review, or must be rejected and regenerated. Bands are deliberately wide so that small reviewer disagreements at the rubric level rarely change the operational decision.

Band	Score range	Meaning	When the output can be used	Required action
Very High Trust	90–100	Output meets the modular framework's	May be used as a finished deliverable	Spot-check 3 citations and the

Band	Score range	Meaning	When the output can be used	Required action
		full quality bar; minor, non-material gaps only.	with sample-based review.	Scorecard. No remediation required.
High Trust	75–89	Output is reliable; one or two categories carry visible weaknesses that do not change the conclusions.	May be used after light human review focused on the weakest category.	Reviewer fixes specific gaps (e.g., missing citation, missing watchlist item). No regeneration.
Moderate Trust	60–74	Output is partially reliable; meaningful gaps in at least one category that affect downstream use.	Must not be used as-is. Suitable for internal drafting or as an input to human-led analysis.	Targeted regeneration of the weakest category, or substantial human edit, before any external use.
Low Trust	40–59	Output has structural or factual weaknesses that make it unsafe to rely on without rework.	Not suitable for use. May be salvaged for skeleton or framing only.	Full prompt review and regeneration. Investigate whether the prompt template was correctly applied.
Very Low Trust	0–39	Output is unreliable. Likely failures include missing structure, fabricated sources, or material factual errors.	Must not be used in any form. Treat as a failed run.	Reject and regenerate. Escalate to prompt engineering if pattern repeats across runs.

Override rules

Two override rules apply on top of the band table, because some failures are too consequential to be averaged out by strong scores elsewhere:

- Hallucinated source override. If any cited source is fabricated, the output is capped at the Low Trust band (40–59) regardless of total score, and the Category 3 score is set to 0.
- Restriction breach override. If a prompt-stated restriction is openly violated, the output is capped at Moderate Trust (60–74) regardless of total score, and Category 1 is reduced by at least one rubric level.

7. Evaluation Checklist

This is the single consolidated checklist reviewers should complete during scoring. It is designed to be used inline — one row per evaluation item, with reviewer notes, observed evidence, the score awarded, and the action required if the score is below threshold.

Evaluation item	Reviewer notes	Evidence observed	Score (0–5)	Action required
1.1 All seven Variables visibly respected				If <4: list which Variables were not respected
1.2 Coverage Period and currency consistent throughout				If <4: flag specific drift points
1.3 Audience and tone match Purpose & Audience field				If <3: rewrite affected sections
2.1 All seven Module 3 sections present and correctly ordered				If <4: list missing/misplaced sections
2.2 Key Metrics Scorecard rendered as a table				If <4: convert to table format
2.3 Module 2 Steps 3–6 visibly addressed in Deep Dive				If <4: identify missing pillar(s)
3.1 Every quantitative claim sourced				If <4: flag unsourced claims for verification

Evaluation item	Reviewer notes	Evidence observed	Score (0–5)	Action required
3.2 Spot-checked citations exist and contain the cited data				If any fail: trigger hallucination override
3.3 Calculations transparent (numerator/denominator shown)				If <3: request derivation rebuild
3.4 Numbers reconcile across Scorecard, Deep Dive, Bottom Line				If <4: identify reconciliation breaks
4.1 Each pillar reaches a defended conclusion				If <4: list descriptive-only pillars
4.2 Peer Context is interpreted, not listed				If <3: request interpretation pass
4.3 Synthesis reconciles tensions across pillars				If <3: regenerate Synthesis section
4.4 Calibrated language used appropriately				If <4: request hedge/confidence rebalance
5.1 Data points dated; staleness flagged where present				If <4: add date stamps
5.2 Material assumptions stated up front				If <3: surface assumptions to opening
5.3 Red flags surfaced in Risks & Red Flags section				If <4: elevate buried red flags

Evaluation item	Reviewer notes	Evidence observed	Score (0–5)	Action required
5.4 All prompt restrictions honoured				If breached: trigger restriction override

Reviewer roll-up: Each block (1.x, 2.x, 3.x, 4.x, 5.x) feeds the corresponding category score. The category score is the average of its rows, rounded to the nearest rubric level (0–5), then converted via the multiplier table in Section 4 to its weighted contribution.

8. Worked Example

The example below applies the framework to a hypothetical financial-health output for a fictional company, NorthArc Industrials plc, FY2024, generated through the modular prompting template described in the source PDF. The example illustrates how rubric levels translate into weighted points and how the band override rules apply.

Input prompt summary

Variable	Value supplied
Company	NorthArc Industrials plc
Identifier	LSE: NARC
Coverage Period	FY2024 (year ended 31 December 2024)
Peer Set	Peer A plc, Peer B SA, Peer C Inc.
Purpose & Audience	Internal credit committee — focused on liquidity and leverage
Data Sources	FY2024 10-K, Q4 2024 trading update, Bloomberg consensus
Accounting/Currency Notes	IFRS; figures in GBP

Reviewer observations

- All seven Module 3 sections were present, in the correct order; the Scorecard was rendered as a table.

- Peer Set was respected with one extension: a fourth peer (Peer D AG) was added without disclosure.
- Most quantitative claims were sourced to the FY2024 10-K; two ratios in the Leverage pillar were stated without showing numerator/denominator.
- Cash flow pillar identified that FCF/net income had fallen from 0.92 to 0.61 over two years, but this red flag was not elevated to the Risks & Red Flags section.
- Forward-looking language in the Catalysts section was appropriately calibrated; the Watchlist contained three specific, monitorable items for the next quarter.
- No hallucinated sources detected on three-source spot-check.

Category-by-category scoring

Category	Max	Rubric level	Multiplier	Points	Justification
1. Prompt Adherence and Instruction Compliance	20	4 (Strong)	0.85×20	17.0	All seven Variables addressed; one peer added beyond the supplied Peer Set; currency consistent.
2. Structural and Modular Framework Alignment	20	5 (Excellent)	1.00×20	20.0	All seven Module 3 sections present, correctly ordered; Scorecard is a table; Steps 3–6 visible.
3. Accuracy, Evidence, and Factual Reliability	25	3 (Adequate)	0.65×25	16.3	Most figures sourced to the FY2024 10-K; two ratios shown without inputs; one peer figure unsourced; no fabrication detected.
4. Reasoning, Judgment, and Evaluation Quality	20	4 (Strong)	0.85×20	17.0	Pillars reach defended conclusions; Synthesis reconciles tensions; Watchlist is specific; Peer Context

Category	Max	Rubric level	Multiplier	Points	Justification
					slightly listed rather than interpreted.
5. Currency, Risk Awareness, Restrictions, and Transparency	15	3 (Adequate)	0.65×15	9.8	Headline figures dated; peer figures undated; one assumption explicit, two implicit; red flag mentioned in Cash flow pillar but not elevated to Risks.

Total Trust Score: $17.0 + 20.0 + 16.3 + 17.0 + 9.8 = 80.1 \rightarrow 80 / 100$

Band: High Trust (75–89).

Override check: No hallucinated source; no restriction breach. No override applies.

Recommended action

Review verdict

Use after light human review. Two specific edits required before the output is sent to the credit committee: (1) elevate the FCF/net income deterioration from the Cash flow pillar into the Risks & Red Flags section; and (2) supply the numerator/denominator for the two leverage ratios. Reviewer should also note the addition of Peer D AG and either remove it or document why it was added.

9. Reviewer Guidance

This section gives reviewers a consistent way to handle the situations that most often cause score drift between reviewers.

Situation	How to handle it
Uncertainty about a category score	Pick the lower rubric level when in doubt, and document the reason in reviewer notes. Consistency is more valuable than charity.

Situation	How to handle it
Missing evidence	Treat absence of evidence as evidence of absence for scoring purposes. Do not assume a citation 'must be there somewhere'. Score against what the reviewer can actually see in the output.
Unsupported claim	Score under Category 3, not Category 4. The deduction belongs to factual reliability, even if the unsupported claim is part of an analytical conclusion.
Outdated information	If the data is older than the Coverage Period implies and is not flagged as such, deduct in Category 5. If it is flagged, no deduction — the transparency is the point.
Hallucination (fabricated source or number)	Apply the hallucinated source override (Section 6). Category 3 is set to 0 and the output is capped at Low Trust regardless of other scores. Document the hallucinated reference verbatim in reviewer notes.
Non-compliance with prompt restrictions	Apply the restriction breach override. Category 1 reduced by at least one rubric level; total score capped at Moderate Trust.
Poor formatting	If formatting issues are cosmetic (font, spacing), do not deduct. If they affect comparability or comprehension (Scorecard rendered as prose, missing section headings), deduct in Category 2.
Weak reasoning despite strong facts	Deduct in Category 4. A factually correct but analytically shallow output is a Category 4 problem, not a Category 3 problem.
Overconfident conclusions	Deduct in Category 4 (calibrated-language criterion). Use Category 5 only if overconfidence is paired with hidden assumptions.
Reviewer disagreement	Escalate to a second reviewer rather than averaging. If the second reviewer's score is in a different band, the lower band applies and the case is logged for the next calibration cycle.

10. Governance and Continuous Improvement

This framework is itself a living artefact and should be governed in step with the modular prompting framework it evaluates. The mechanisms below are recommended and should be adapted to the organisation's broader AI governance regime.

Note: the cadences and mechanisms in this section are professional assumptions, not derived from the source PDF. The PDF describes the modular prompting framework but does not prescribe its own governance.

Ownership

- A single named owner — typically AI Governance or QA leadership — owns the framework's content, version history, and calibration log.
- Prompt engineering owns the modular prompt template (Variables, Module 1, Module 2, Module 3); coordination between the two owners is required whenever either changes.

Calibration cadence

- Quarterly: a calibration session in which two reviewers independently score the same five recent outputs. Score gaps greater than one band trigger a guidance update.
- Ad-hoc: any time the modular prompt template adds, removes, or renames a Variable or a Module section, the framework is reviewed to confirm category criteria still align.

Versioning

- The framework is versioned (v1.0, v1.1, v2.0). Major versions reflect category changes; minor versions reflect rubric or guidance refinement.
- Each output's Trust Score should record the framework version under which it was scored, so historical comparisons remain meaningful.

Continuous improvement signals

- Recurring failure patterns observed in reviewer notes feed back into the prompt template — typically into Module 1 (Research Standard) or Module 3 (Output Structure).
- Rubric levels that produce frequent reviewer disagreement are candidates for sharper criteria or worked examples.
- Override rules are reviewed annually; new override candidates require evidence of repeated, high-consequence failure.

11. Final Summary

The AI Output Trust Score Evaluation Framework converts the modular prompting architecture in the source PDF into an evaluation contract. Where the prompt template enforces a quality bar, a methodology, and a deliverable format on the input side, this framework enforces the matching standards on the output side, using five weighted categories that together total 100 points.

By scoring Prompt Adherence, Structural Alignment, Accuracy and Evidence, Reasoning Quality, and Currency and Transparency separately and explicitly, the framework gives reviewers a defensible audit trail, gives prompt engineers a feedback loop, and gives business stakeholders a single trust signal that

maps directly to a usage decision. The interpretation bands and override rules ensure that the most consequential failures, fabricated sources, restriction breaches, cannot be averaged away by surface polish elsewhere.

Used consistently, the framework improves three things at once: confidence in individual outputs, accountability across the AI delivery chain, and consistency between reviewers, periods, and use cases. It is intended to evolve alongside the modular prompting template itself, with calibration and versioning as standing practices rather than one-off projects.

Bottom line

Modular prompting fixes how AI outputs are produced. The Trust Score framework fixes how they are evaluated. Together, they make analytical AI deliverables comparable, defensible, and fit for downstream business use.